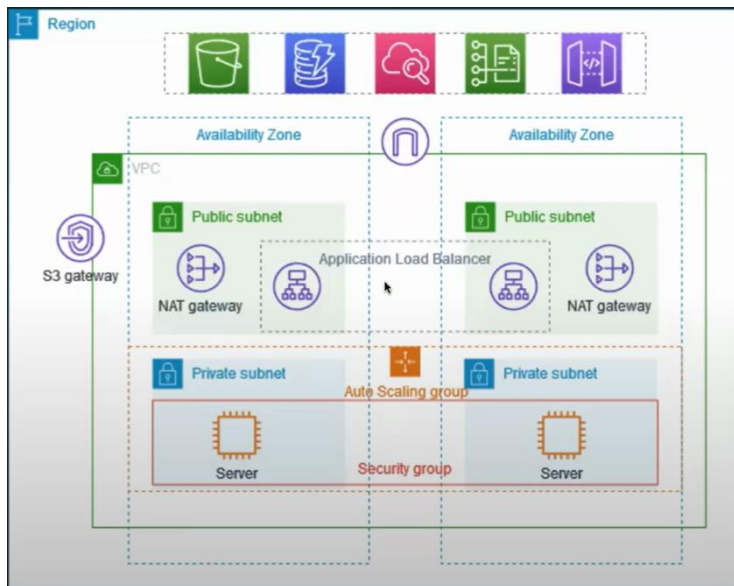


Project Report: Implementing a VPC for a Production Environment

Overview

This project involves creating a Virtual Private Cloud (VPC) to host servers in a production environment. The VPC will provide a secure and scalable network infrastructure within a public cloud, utilizing multiple Availability Zones (AZs) for high availability and redundancy. Key components include an Auto Scaling group, an Application Load Balancer, private and public subnets, and a NAT gateway for internet access.



Terminologies

Virtual Private Cloud (VPC)

A Virtual Private Cloud (VPC) is a private network within a public cloud, such as AWS or Google Cloud, allowing complete control over network settings like IP address ranges, subnets, and security configurations. It provides an isolated section of the cloud to run resources securely with the flexibility and scalability of the cloud.

Network Address Translation (NAT)

Network Address Translation (NAT) maps private IP addresses to a public IP address or a pool of public IP addresses. This enables multiple devices on a private network to access the internet using a single public IP address, enhancing security and conserving IP address space.

Infrastructure

1. Resiliency with Multiple Availability Zones (AZs)

- Servers are deployed across two AZs to ensure redundancy and high availability. Each AZ is a separate data center within the same region.

2. Auto Scaling Group

- Automatically scales the number of servers based on demand to ensure sufficient resources are available.

3. Application Load Balancer

- Distributes incoming traffic across servers to improve availability and performance.

4. Private Subnets

- Servers are deployed in private subnets for security, meaning they are not directly accessible from the internet.

5. NAT Gateway

- Allows servers in private subnets to access the internet for software updates or external services.

6. Resilient NAT Gateway

- Deployed across both AZs to ensure continuous internet connectivity even if one AZ fails.

Project Architecture

The VPC is designed with both public and private subnets across two Availability Zones. Each public subnet contains a NAT gateway and a load balancer node. Servers run in private subnets, launched and terminated by an Auto Scaling group, and receive traffic from the load balancer. The servers can access the internet via the NAT gateway.

Components and Configuration

VPC

- **CIDR Block:** 10.0.0.0/16

Subnets

- **Public Subnets:**
 - Subnet 1 (AZ1): 10.0.1.0/24
 - Subnet 2 (AZ2): 10.0.2.0/24

- **Private Subnets:**

- Subnet 3 (AZ1): 10.0.3.0/24
- Subnet 4 (AZ2): 10.0.4.0/24

Internet Gateway (IGW)

- Attach an Internet Gateway to the VPC to allow public subnets to access the internet.

NAT Gateway

- Deploy a NAT Gateway in each public subnet to provide internet access to private subnets.

Route Tables

- **Public Route Table:** Route 0.0.0.0/0 to IGW.
- **Private Route Table:** Route 0.0.0.0/0 to NAT Gateway.

Auto Scaling Group

- Launches and terminates servers in private subnets based on demand.

Application Load Balancer

- Distributes incoming traffic across servers in private subnets.

Implementation Procedure**Step 1: Create a VPC**

1. Create a VPC with a CIDR block of 10.0.0.0/16.

Step 2: Create Subnets

1. Create two public subnets in different AZs with CIDR blocks 10.0.1.0/24 and 10.0.2.0/24.
2. Create two private subnets in different AZs with CIDR blocks 10.0.3.0/24 and 10.0.4.0/24.

Step 3: Configure Internet Gateway

1. Create an Internet Gateway and attach it to the VPC.
2. Create a public route table and add a route for 0.0.0.0/0 pointing to the Internet Gateway.
3. Associate the public route table with the public subnets.

Step 4: Create NAT Gateway

1. Create a NAT Gateway in each public subnet.

2. Create a private route table and add a route for 0.0.0.0/0 pointing to the NAT Gateway.
3. Associate the private route table with the private subnets.

Step 5: Configure Auto Scaling Group

1. Create a launch configuration or launch template specifying the AMI, instance type, and security group.
2. Create an Auto Scaling group with the desired capacity, minimum and maximum instances, and target the private subnets.

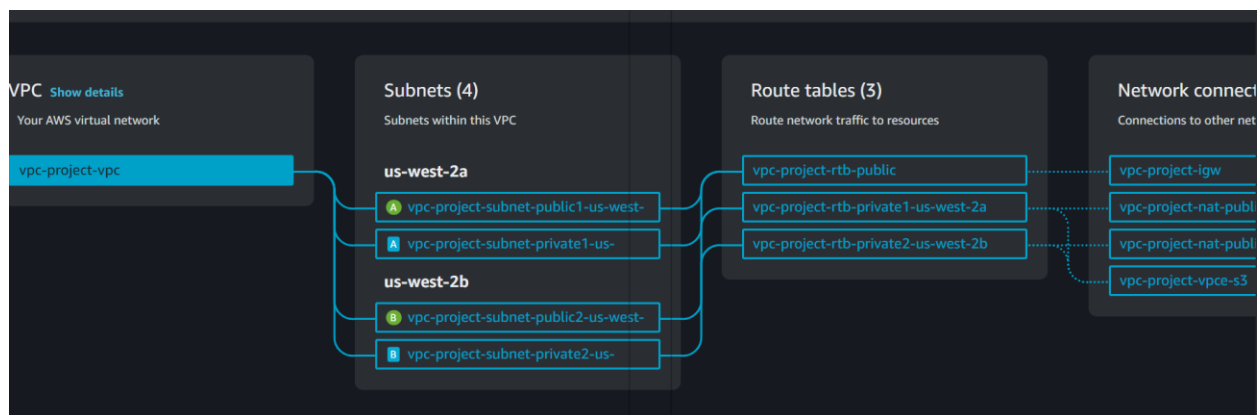
Step 6: Configure Application Load Balancer

1. Create an Application Load Balancer and specify the public subnets.
2. Create target groups and register the Auto Scaling group instances.
3. Create listeners to route traffic from the load balancer to the target groups.

Related Images :

The screenshot shows the 'Create VPC' page in the AWS Management Console. The breadcrumb navigation is 'VPC > Your VPCs > Create VPC'. The page title is 'Create VPC' with an 'Info' link. A descriptive text states: 'A VPC is an isolated portion of the AWS Cloud populated by AWS objects, such as subnets, route tables, and network interfaces.' The 'VPC settings' section includes:

- Resources to create:** Two radio buttons are present: 'VPC only' (unselected) and 'VPC and more' (selected).
- Name tag auto-generation:** A checkbox labeled 'Auto-generate' is checked. Below it, a text input field contains the value 'project'.
- IPv4 CIDR block:** A text input field contains '10.0.0.0/16', and to its right, it indicates '65,536 IPs'. A note below states: 'CIDR block size must be between /16 and /28.'
- IPv6 CIDR block:** Two radio buttons are present: 'No IPv6 CIDR block' (selected) and 'Amazon-provided IPv6 CIDR block' (unselected).



VPC > Your VPCs > Create VPC > Create VPC resources

Create VPC workflow

✔ Success

▼ Details

✔ Create VPC: vpc-04280eae804d70bb8

✔ Enable DNS hostnames

✔ Enable DNS resolution

✔ Verifying VPC creation: vpc-04280eae804d70bb8

✔ Create S3 endpoint: vpce-06db1eadd4e4b5676

✔ Create subnet: subnet-0a15383ec60b24d68

✔ Create subnet: subnet-02811e9ade8465de5

✔ Create subnet: subnet-0966da2aa5afc0dbd

✔ Create subnet: subnet-0b681ec6b64c63069

✔ Create internet gateway: igw-0c5ec146628cf5f6a

✔ Attach internet gateway to the VPC

✔ Create route table: rtb-047b43c819b03a0be

✔ Create route

✔ Associate route table

✔ Associate route table

✔ Allocate elastic IP: eipalloc-0b27313af793785b4

✔ Allocate elastic IP: eipalloc-0372ae1b1f298e33c

✔ Create NAT gateway: nat-0ba1ad6854cec589c

VPC > Your VPCs > vpc-04280eae804d70bb8

vpc-04280eae804d70bb8 / vpc-project-vpc

Actions ▼

DetailsInfo

VPC ID vpc-04280eae804d70bb8	State ✔ Available	DNS hostnames Enabled	DNS resolution Enabled
Tenancy Default	DHCP option set dopt-07469e66506c63099	Main route table rtb-064f3186c4d5efea3	Main network ACL acl-0cb63633c5b8fbd8c
Default VPC No	IPv4 CIDR 10.0.0/16	IPv6 pool -	IPv6 CIDR (Network border group) -
Network Address Usage metrics Disabled	Route 53 Resolver DNS Firewall rule groups -	Owner ID 891376971929	


Step 1: Choose launch template

[Edit](#)

Group details

Auto Scaling group name
myautoscalinggroup

Launch template


Launch template	Version	Description
aditiautoscaling  lt-050cc2b6b34b9b759	Default	A auto scaling group



Step 2: Choose instance launch options

[Edit](#)

Network

Network

VPC
[vpc-04280eae804d70bb8](#) 

Availability Zone	Subnet	
us-west-2b	subnet-0b681ec6b64c63069 	10.0.144.0/20
us-west-2a	subnet-0966da2aa5afc0dbd 	10.0.128.0/20

Group size

Desired capacity	Desired capacity type
2	Units (number of instances)

Scaling

Minimum desired capacity	Maximum desired capacity
1	4
Target tracking policy	-

EC2 > Auto Scaling groups

Auto Scaling groups (1/1) Info Refresh Launch configurations Launch templates Actions Create Auto Scaling group

<input checked="" type="checkbox"/>	Name	Launch template/configuration	Instances	Status	Desired capacity
<input checked="" type="checkbox"/>	myautoscalinggroup	aditiautoscaling Version Default	2	-	2

Instances (2) Info Refresh Connect Instance state Actions Launch instances

All states

Instance state = running Clear filters

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status
<input type="checkbox"/>		i-0974869ff3c6734ed	Running	t2.micro	Initializing	View alarms
<input type="checkbox"/>		i-075e9fc8e9e97794f	Running	t2.micro	Initializing	View alarms
<input checked="" type="checkbox"/>	BastionHost	i-0db306061e4a5f29b	Running	t2.micro	Initializing	View alarms

A public server will be intermediate between private servers and the internet


```
aditi@ADITI MINGW64 ~/downloads
$ ssh -i aditi.pem ec2-user@34.222.236.145
Last login: Fri May 24 18:25:56 2024 from 49.34.209.43

#_
~\#### Amazon Linux 2
~~\#####
~~\###| AL2 End of Life is 2025-06-30.
~~\#/
~~V~'-'>
~~~~
~~~~
~/m/'-

A newer version of Amazon Linux is available!

Amazon Linux 2023, GA and supported until 2028-03-15.
https://aws.amazon.com/linux/amazon-linux-2023/

[ec2-user@ip-10-0-19-197 ~]$ ls
aditi.pem
[ec2-user@ip-10-0-19-197 ~]$ ssh -i aditi.pem ec2-user@10.0.157.228
The authenticity of host '10.0.157.228 (10.0.157.228)' can't be established.
ECDSA key fingerprint is SHA256:s6SRsQppv3Sig46mnxvcQl3o43gYftQKUSlP2xDfxQ.
ECDSA key fingerprint is MD5:ba:19:62:2c:aa:7d:00:30:b9:d7:93:85:8a:6e:c4:12.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '10.0.157.228' (ECDSA) to the list of known hosts.
@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@
@                WARNING: UNPROTECTED PRIVATE KEY FILE!                @
@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@
Permissions 0444 for 'aditi.pem' are too open.
It is required that your private key files are NOT accessible by others.
This private key will be ignored.
Load key "aditi.pem": bad permissions
Permission denied (publickey,gssapi-keyex,gssapi-with-mic).
[ec2-user@ip-10-0-19-197 ~]$ |
```

```
ec2-user@ip-10-0-19-197 ~]$ vim index.html
ec2-user@ip-10-0-19-197 ~]$ vim aditi.html
ec2-user@ip-10-0-19-197 ~]$
```

```
[ec2-user@ip-10-0-19-197 ~]$ python3 -m http.server 8000
Serving HTTP on 0.0.0.0 port 8000 (http://0.0.0.0:8000/) ...
```

myloadbalancer

Actions

▼ Details

Load balancer type

Application

Scheme

Internet-facing

Status

Provisioning

Hosted zone

Z1H1FL5HABSF5

VPC

vpc-

04280eae804d70bb8

Availability Zones

subnet-

0a15383ec60b24d68

us-west-2a (usw2-az1)

subnet-

02811e9ade8465de5

us-west-2b (usw2-az2)

IP address type

IPv4

Date created

May 25, 2024, 00:30 (UTC+05:30)

Load balancer ARN

arn:aws:elasticloadbalancing:us-west-2:891376971929:loadbalancer/app/myloadbalancer/fabc1efbe7b8a3c4

DNS name

Info

myloadbalancer-863426570.us-west-2.elb.amazonaws.com (A Record)

<

Listeners and rules

Network mapping

Resource map - new

Security

Monitoring

Integrations

>

Listeners and rules (1) Info

Manage rules

Manage listener

Add listener

A listener checks for connection requests on its configured protocol and port. Traffic received by the listener is routed according to the default action and any additional rules.

Q Filter listeners

< 1 >

Protocol:Port

▼

Default action

▼

Rules

▼

ARN

▼

Security

▼

HTTP:80

Forward to target group

• mygrouptarget: 1 (100%)

• Group-level stickiness: Off

1 rule

ARN

Not appli



Conclusion

By following this procedure, you can implement a robust and scalable VPC architecture for a production environment. The architecture ensures high availability, security, and efficient resource utilization through the use of multiple Availability Zones, Auto Scaling, and load balancing. The NAT gateway provides secure internet access for servers in private subnets, completing a well-rounded infrastructure setup.